

Object-focused Storage and Management of Science Data


Suren Byna

Lawrence Berkeley National Laboratory

sbyna@lbl.gov

FlexScience 2022

July 1st, 2022




Credits


- Contributions from:

- PDC (Proactive Data Containers) project team
 - Houjun Tang, Kimmy Mu, Jerome Soumagne, Richard Warren, Teng Wang, Bin Dong, François Tessier, Quincey Koziol, Qiao Kang, Marc Snir, Chen Wang
- ExaHDF5 / ExaIO team
 - Suren Byna, Scot Breitenfeld, Venkat Vishwanath, Houjun Tang, Qiao Kang, Jean Luca Bez, Huihuo Zheng, John Mainzer, Quincey Koziol*, Neil Fortner, Dana Robinson, Jordan Henderson, Jerome Soumagne, Richard Warren, John Mainzer, Neelam Bagha, Elena Pourmal, Michela Becchi, John Ravi, Wei Zhang, Yong Chen, Kathryn Mohror, Sarp Oral, Adam Moody, Cameron Stanavige, Michael Brim, Seung-Hwan Lim, Ross Miller, Swen Boehm



Scientific data storage and access - Current scenario







Outline of this talk (focusing on HPC systems)

- Basics, trends, challenges, and requirements
- Object-focused data management
- Optimizations in HDF5
- Proactive Data Containers (PDC) data management runtime system
- Future research directions

I/O software stack - Several layers with inter-dependencies



Architectural trends impacting I/O - deep memory and storage hierarchy



Trends in computing devices

- Heterogeneous processing devices
 - CPUs
 - GPUs
 - FPGAs
 - Special purpose accelerators
 - ...
- Massive concurrency
- Locations of data generation and consumption
 - Traditional: In compute nodes
 - Trends: In network, in storage, and at the edge






Image from D. Vasudevan, via J. Shalf,
Extreme Heterogeneity workshop report



Application I/O Trends

- HPC simulation and analysis
 - Parallel applications storing or retrieving data
 - checkpoint, restart, and analysis - mostly defensive I/O
 - Partitioned regions of data in a large data object
 - Each process accesses (writes or reads) a single or multiple regions
 - Typically non-overlapping (there may be some overlapping at region boundaries)
 - Number of files: Single shared file ($N-1$), file-per-process ($N-N$), a few sub-files ($M-N$)
 - APIs: POSIX, MPI-IO, HDF5, etc.



Application I/O Trends

- HPC simulation and analysis
 - Parallel applications storing or retrieving data
 - checkpoint, restart, and analysis - mostly defensive I/O
 - Partitioned regions of data in a large data object
 - Each process accesses (writes or reads) a single or multiple regions
 - Typically non-overlapping (there may be some overlapping at region boundaries)
 - Number of files: Single shared file ($N-1$), file-per-process ($N-N$), a few sub-files ($M-N$)
 - APIs: POSIX, MPI-IO, HDF5, etc.
- ML / AI
 - Read-heavy
 - Random accesses to different parts of a file or multiple files (due to shuffling between epochs)
 - Small I/O requests to batches of data, metadata accesses
 - Large numbers of files - $O(100,000)$
 - APIs: Python, ML/AI frameworks (TensorFlow, PyTorch, Caffe, etc.)
- Experimental and observational facilities, Edge computing, etc.
 - Large numbers of small files, streaming data, sparse data, remote file accesses



What do users want?

Use case	Domain	Sim/EOD/analysis	Data size	I/O Requirements
FLASH	High-energy density physics	Simulation	~1PB	Data transformations, scalable I/O interfaces, correlation among
CMB / Planck	Cosmology	Simulation,	10PB	Automatic data movement
Easy interfaces to complex systems				
				transformations
Autonomous data movement and performance tuning				
Information capture, management, and search				
TECA	Climate	Analysis	~10PB	Data organization and efficient data movement
HipMer	Genomics	EOD/Analysis	~100TB	Scalable I/O interfaces, efficient and automatic data movement

Users:


- Scientists
- App developers
- Supercomputing facilities
- System designers



Storage and I/O challenges

- Poor performance
 - POSIX semantics
 - Tuning options at different software layers
- Complexity with deep storage hierarchy
 - Using the performance and capacity storage layers efficiently and transparently
- Heterogeneity of storage devices
 - From memory-class to hard-disks to tape
- Variability of performance on HPC systems
 - Some part of the storage system is concurrently shared by multiple users
- New classes of applications
 - ML / AI applications are read-heavy with random access patterns
 - EOD and Edge computing applications
- New high-level programming models
 - Performance **AND** Productivity

Object-focused data management to the rescue?



Object storage

- Designed for unstructured data such as media, documents, logs, backups, application binaries and VM images
- Data objects are associated with metadata descriptions
- Common API: REST
- Examples: Ceph, Swift, Amazon S3, etc.








Objects - An overloaded term

Object management system	What does object mean?
Parallel file systems (Lustre, etc.)	Chunks of a file (block storage)
Cloud object storage (S3, etc.) OpenStack Swift, MarFS, Ceph, etc.	Files (images, videos, etc.) + metadata
HDF5 datasets	Multi-dimensional arrays, images, any type of data
DAOS	Multi-dimensional arrays, files

HDF5 self-describing file format and API for science apps


- HDF5 is a self-describing file format, API, and tools designed to store, access, analyze, share, and preserve diverse, complex data in continuously evolving heterogeneous computing and storage environments
 - Maintained by The HDF Group (THG)
- Heavily used on DOE supercomputing systems and diverse science applications across the world
- Many ECP applications have dependency on HDF5-based I/O
 - 17 critical
 - 11 important
 - 8 interested



HDF5 Containers (Files), and Data / Metadata objects

HDF5 files, groups,
and links
organize
data objects.

HDF5 datasets
and attributes
store
application data.





HDF5 Attributes

- Typically contain user metadata
- Have a name and a value
- Attributes “decorate” HDF5 objects
- Value is described by a datatype and a dataspace
- Analogous to a dataset, but do not support partial I/O operations
 - Nor can they be compressed or extended



HDF5 Home Page

HDF5 home page: <http://www.hdfgroup.org/solutions/hdf5/>

- Latest releases: HDF5 1.12.1 and 1.13.1


HDF5 source code:

- Written in C, and includes optional C++, Fortran, and Java APIs
 - Along with “High Level” APIs
- Contains command-line utilities (h5dump, h5repack, h5diff, ..) and compile scripts
- <https://github.com/HDFGroup/hdf5>


HDF5 pre-built binaries:

- When possible, include C, C++, Fortran, Java and High Level libraries.
 - Check ./lib/libhdf5.settings file.
- Built with and require the SZIP and ZLIB external libraries

ExaHDF5 - Our enhancements to HDF5 and beyond



Massive concurrency, abundant computing power



Deep hierarchy of memory and storage

Easy interfaces to complex systems

Autonomous data movement and performance tuning

Information capture, management, and search

Goals




Our enhancements to HDF5 and beyond

- Asynchronous I/O
- Caching
- GPU I/O
- Subfiling (in development)
- Multi-dataset I/O (in development)
- Data reduction with parallel compression
- Understanding parallel I/O performance
- Autotuning parallel I/O
- Performance tuning for applications
- Next generation data management systems


Asynchronous I/O

- HPC simulations and several analysis have iterative I/O phases
 - Computation phase and I/O phase
 - I/O phase is typically to perform checkpointing or storing the current state of the simulation
- Asynchronous I/O: Overlap I/O phase with compute phase




Asynchronous I/O

- HPC simulations and several analysis have iterative I/O phases
 - Computation phase and I/O phase
 - I/O phase is typically to perform checkpointing or storing the current state of the simulation
- Asynchronous I/O: Overlap I/O phase with compute phase




Asynchronous I/O

- HPC simulations and several analysis have iterative I/O phases
 - Computation phase and I/O phase
 - I/O phase is typically to perform checkpointing or storing the current state of the simulation
- Asynchronous I/O: Overlap I/O phase with compute phase




HDF5 Async I/O Implementation

- Using background threads
- Asynchronous task queue
- Uses an “event set” to manage async operations
- Transparent background thread execution



Transparent Task Dependency Management

- All I/O operations can only be executed after a successful file create/open.
- A file close operation can only be executed after all previous operations in the file have been completed.
- All read or write operations must be executed after a prior write operation to the same object.
- All write operations must be executed after a prior read operation to the same object.
- All collective operations must be executed in the same order with regard to other collective operations.
- Only one collective operation may be in execution at any time.




Challenge: How to perform asynchronous I/O transparently from the user?

Async I/O – Low overhead and efficient hiding of I/O latency


AMReX: A software framework for massively parallel, block-structured adaptive mesh refinement (AMR) applications

AMReX Single-level Plotfile 385GB x 5 timestep on Summit



Nyx: An adaptive mesh, cosmological hydrodynamics simulation code (left)

AMReX Multi-level Plotfile 559GB x 5 timesteps on Summit




Castro: An adaptive mesh, astrophysical radiation/MHD/hydrodynamics simulation code

Async I/O VOL Connector

- Available now:
 - Source: <https://github.com/hpc-io/vol-async>
 - Docs: <https://hdf5-vol-async.readthedocs.io/en/latest>
- Future work:
 - Merge compatible VOL operations
 - If two async dataset write operations are putting data into same dataset, can merge into only one call to underlying VOL connector
 - Turn multiple ‘normal’ group create operations into a single ‘multi’ group create operation
 - Use multiple background threads
 - Needs HDF5 library thread-safety work, to drop global mutex
 - Switch to TaskWorks thread engine
 - A portable, high-level, task engine designed for HPC workloads
 - Task dependency management, background thread execution.

Cache VOL Connector - Integrating node-local storage into parallel I/O





Theta @ ALCF: Lustre + SSD (128 GB / node),
ThetaGPU (DGX-3) @ ALCF: NVMe (15.4 TB / node)
Summit @ OLCF: GPFS + NVMe (1.6 TB / node)

Cache VOL

- Using node-local storage for caching / staging data for fast and scalable I/O.
- Data migration to and from the remote storage is performed in the background.
- Managing data movement in multi-tiered memory / storage through stacking multiple VOL connectors (*async -> cache -> async*)
- All complexity is hidden from the users

Repo: <https://github.com/hpc-io/vol-cache>


Parallel Write (H5Dwrite) with cache VOL



Details are hidden from the application developers.

Parallel Read (H5Dread) w/ cache VOL


Create memory mapped files and attached them to a MPI_Win for one-sided remote access



First time reading the data


One-sided communication for accessing remote node storage.

- Each process exposes a part of its memory to other processes (MPI Window)
- Other processes can directly read from or write to this memory, without requiring that the remote process synchronize (MPI_Put, MPI_Get)




Reading the data directly from node-local storage

HDF5 Cache VOL uses node-local SSDs as cache → 2X to 3X faster



VPIC-IO is a kernel derived from a plasma physics simulation of solar weather interacting with the earth's magnetosphere. The simulation writes particle data to HDF5 file, where each variable is mapped to a HDF5 dataset



CosmoFlow is a 3D convolutional neural network model for learning the universe at scale. The model is implemented in TensorFlow with Horovod for data parallel training at scale.


What do users want?

Use case	Domain	Sim/EOD/analysis	Data size	I/O Requirements
FLASH	High-energy density physics	Simulation	~1PB	Data transformations, scalable I/O interfaces, correlation among
CMB / Planck	Cosmology	Simulation,	10PB	Automatic data movement
Easy interfaces to complex systems				
				transformations
Autonomous data movement and performance tuning				
Information capture, management, and search				
TECA	Climate	Analysis	~10PB	Data organization and efficient data movement
HipMer	Genomics	EOD/Analysis	~100TB	Scalable I/O interfaces, efficient and automatic data movement

Users:


- Scientists
- App developers
- Supercomputing facilities
- System designers

Current file-based approach is still burdensome




- **Challenges**
 - Multi-level hierarchy complicates data movement, especially if user has to be involved
 - POSIX-IO semantics hinder scalability and performance of file systems and IO software

Object-centric data management





Proactive Data Containers (PDC) - An object-focused data movement runtime system - Abstractions

- Container – special case of object, metadata only
 - For grouping objects
- Object - metadata and data payload
 - Multi-dimensional arrays
- Regions
 - A logical region in a multi-dimensional array
 - expressed with offsets, sizes, and strides
 - multi-dimensional



Proactive Data Containers (PDC) - An object-focused data movement runtime system



- Advantages with PDC
 - Application-level object abstractions - Freedom from file management
 - Transparent utilization of storage hierarchy and data movement
 - Superior and scalable performance
 - Live system for data management services
 - Metadata management, analysis, indexing and querying services, consistency, data placement, etc.




PDC Services

- Metadata management - to store and query objects' metadata
- Data management - to move data efficiently and asynchronously using multi-tier storage
- Data transformation and analysis - Perform data transformations or analysis while data is in flight

PDC metadata management


- Metadata is critical for finding objects previously written
- Predefined tags must be filled - either by users (using API) or by extracting from system
- User-defined tags improve findability of objects
 - KV pairs
- Unique IDs for metadata objects
- Querying for objects using tags
 - Exact match
 - Partial match

Pre-defined Tag	User-defined Tag
<ul style="list-style-type: none">• Object ID• Data Location• System Info• ID Attributes<ul style="list-style-type: none">- Object Name- Application name- Ownership- Timestep	<ul style="list-style-type: none">• (Tag Name0, Value0)• (Tag Name1, Value1)• (Tag Name2, Value2)• ...




Object-centric PDC API

- No explicit data movement
- Container
 - create container
 - delete container
 - add / delete objects
- Objects
 - create object
 - add metadata
 - create regions
 - map objects / regions from source to destination
 - Source and destinations can be memory or PDC spaces
 - lock when updating an object
 - release informs PDC runtime for implicit data movement
 - find object (followed by “map” for reading)
 - Explicit put and get object functions are also available



Transparent data movement in storage hierarchy

- Usage of compute resources for I/O
 - Shared mode – Compute nodes are shared between applications and I/O services
 - Dedicated mode – I/O services on separate nodes
- Users start PDC server and application links with the PDC client library
 - Set environment variables for informing PDC about the memory and storage locations
- Transparent data movement by PDC servers
 - Apps map data buffers to objects and PDC servers place and manage data
 - Apps query for data objects using attributes






Experimental setup


- Cori at NERSC
 - 1600 “Haswell” compute nodes
 - 128 GB DRAM/node
 - 32 cores/node
- SSD-based “Burst Buffer”
- HDD-based shared file system Lustre
 - *Shared* mode: One PDC server on each node, the remaining 31cores for user application execution.
 - *Dedicated* mode: PDC servers and user’s application are on separate nodes
- Mercury Remote Procedure Calls (RPCs), as the communication mechanism between client and server and between servers - TCP and Cray GNI
- Benchmarks
 - VPIC-IO: I/O of a large-scale plasma physics simulation code
 - BD-CATS-IO: I/O kernel of a big data clustering analysis code

VPIC-IO (Weak Scaling) Multi-timestep Write



Total time to write 5 timesteps from the VPIC-IO kernel to Lustre and Burst Buffer on Cori. PDC is up to **5x** faster than HDF5 and **23x** over PLFS.


BD-CATS-IO (Weak scaling) Multi-timestep Read



Total time for reading data of 5 timesteps from the BD-CATS-IO kernel from the Lustre and from the burst buffer. PDC is up to **11X** faster than PLFS and HDF5.

PDC analysis service - Traditional analysis

- Traditional data analysis - store data in disks or tapes and access it later for analysis (high latency in moving data)
- In-situ analysis reduces data access latency, but coordination between simulations and analysis codes is required
- User-Defined Functions (UDF) reduces user involvement - current solutions are designed for post-processing



- Requirements for transparent analysis strategy:
 - Reduce user involvement and data production and consumption stages
 - Analyze data transparently while the data is in movement

PDC analysis service - “In-locus” analysis




PDC API to register analysis functions and transforms

Analysis Registration:

```
PDCobj_analysis_register("user-defined-analysis-function", input1_iter,  
result1_iter);
```


Transform Registration:

```
PDCbuf_map_transform_register("pdc_transform_compress", &x[0], region_x,  
obj_xx, region_xx, 0, INCR_STATE, DATA_OUT);
```




PDC VOL Connector for HDF5 Applications

- Currently only implements a subset of the HDF5 API
 - HDF5 files mapped to PDC containers
 - HDF5 datasets mapped to PDC objects
 - PDC regions similar to HDF5 selections
- File create, open, and close operations are a direct match to PDC container operations



Current status of PDC



- Current achievements in a next generation intelligent data management framework
 - Transparent
 - Asynchronous
 - High performance and scalable
 - Partial autonomy in utilizing storage hierarchy


Source code:

<https://github.com/hpc-io/pdc>

Publications


<https://sdm.lbl.gov/pdc/>

Proactive Data Containers (PDC) - Ongoing and future work




- Interface with Python and object (put and get) interfaces
- Extend to KV stores abstractions
- Data reorganization and placement based on access history
- Dynamic consistency
- Byte address mapping with PDC storage objects - working with John Shalf et al.
- Analysis on storage and network computing resources
- Federated data objects - remote access across multiple file systems at a site, multiple HPC sites, and cloud environments

Scientific data storage and access - Looking into future



Future research to achieve autonomous, object-focused data management





Conclusions

- Trends
 - Architecture - deepening hierarchy, high concurrency
 - Applications - increasing diversity (exascale, EOD, ML/AI, etc.)
 - User requirements - productivity and performance, knowledge management
- Our recent solutions at the I/O software level
 - Asynchronous I/O
 - Caching
 - Object-centric data management runtime system
- **Domain scientists should not be burdened with data (I/O and other data movement) and metadata management tasks**
 - Achieving autonomous storage and I/O requires novel high-productivity interfaces, information extraction and management, automatic data movement, reconfigurable storage, etc.